

Enterprise Architecture & Governance

12-month engagement

Prepared for Ty Staviski, Deputy CFO, Techtronic Industries

Julian Hart · August 2026 · Private & confidential

At a glance: 12 months, US\$299,563 all-in, breaking even on a recurring saving of US\$132,500/year.

I'd work on the initiatives you already have, in an enterprise-architecture context. Not a parallel programme, not a new function competing for sponsorship.

The North American ERP conversion is proof of what's possible: a hard project done right because one team rallied around it. The risk with a win like that in a federated business is the lessons stay with the people who lived it and never reach the next region that needs them. Standards rarely fail at approval. They fail in adoption, quietly, region by region, until the reconciliation comes back.

You named the levers yourself: leverage synergies, AI and your assets. We'd build the governed layer that puts those levers to work, against your own target of 10% EBIT. With SG&A rising 173 bps in H1, we need to identify additional levers that ensure gross margin gains reach EBIT.

The next point could well come from the structural side. Steve asked exactly that at the annual results briefing: *"How can we leverage our negotiating costs on IT? How can we do the things globally to be able to free up more cash?"* Cost management frameworks like FinOps and TBM enable informed trade-off conversations between IT and finance. Architecture frameworks like TOGAF ensure adoption of global standards. It also feeds the free cash flow target raised mid-year. A region that reuses a group standard reduces capital up front, and robust assurance reduces capital spend on rework.

Where the cost reduction is: five levers

Lever	Mechanism
Capability mapping	In any federated business without a capability map, investment decisions default to the projects with the strongest sponsors, because there is no shared view of which capability each project serves or how it advances group priority. The map provides that view: every active technology project tied to the capability it serves, with a fund, defer or kill recommendation logged for each. The metric is portfolio coverage: the share of active projects mapped and carrying a logged recommendation. The investment decision stays with you; what changes is that it's made against the map, not the sponsor.
Portfolio rationalisation	Every application mapped to the capability it serves, with its run cost against it. Where two regions run separate platforms for the same capability (finance platforms are the classic case), each duplicate carries its own licences, support contracts and upgrade cycle. Rationalising at contract end takes out run cost that recurs every year: a saving on money you are already spending, not a projected benefit.
Decision rights	Cross-region decisions escalate to committee because no one holds the right to make them. The governed layer gives each class of decision a named owner, a standard to decide against, a ratifier and a clock. The metric is time-to-decision: month one baselines how long cross-region calls take to clear today, and that number, tracked monthly, is what you measure me against. Every week a decision stops waiting is a week its benefit stream stops being deferred.
Design assurance	Programmes that self-certify their own design are how overruns and late re-cuts happen: the people accountable for the date are also the people judging whether the design is sound. An independent assurance gate separates those two jobs, so problems surface while they're still cheap to fix. The metric is coverage: the share of major programmes clearing independent

design review before their dates lock. Overruns are your number, not mine. But rework falls where assurance holds, and the NA ERP conversion shows what that discipline looks like when it's in the room.

Vendor standardisation Steve's point on the results call: leveraging negotiating costs on IT to free up cash. A reference architecture defines the standard platform per capability, so the estate converges from many regional variants toward few group standards. The metric is variant count per capability: baselined, with reduction candidates logged against contract end dates. Fewer variants means fewer, larger negotiations: the leverage Steve asked for. The negotiated saving is procurement's number; the variant count that makes it possible is mine.

Scope: three capabilities, one engagement

All of it inside the enterprise-architecture context, working on the initiatives you already have. These are not three programmes; they are three applications of the same architecture layer: standards, decision rights, cost visibility and reusable patterns.

1. The architecture capability

Standing it up follows TOGAF's Preliminary Phase, tailored to TTI rather than adopted wholesale:

- **The capability itself:** which parts of the enterprise it covers, at what depth, and the target maturity: what you need now, evolving from there, not a full-dress framework on day one.
- **The operating model:** where the architecture function sits, the roles around it, and who is accountable.
- **Governance:** decision rights, approval and escalation paths, and how exceptions are handled.
- **Principles:** the enterprise-wide guardrails future architecture decisions get made against.
- **Fit with what you already run:** the frameworks, methods and delivery practices in place today; the capability is designed around them, not dropped on top of them.
- **Tools and skills:** the architecture repository, and the competency gaps that need closing for TTI's own teams to operate the capability.

The design has to be group-level and the effort concentrates wherever the alignment gap actually is. Group design, targeted effort, right-sized maturity: a group capability without paying for coverage you don't need.

2. Cost transparency across the technology estate

You can't reduce a cost estate you can't read, and in most federated groups technology sits as one lump in the P&L. Two disciplines exist for this: TBM is the strategic, top-down version; FinOps is the real-time, operational one. Same principle: see cost clearly, tie it to value, decide together. A full TBM implementation is heavy machinery (agreeing the taxonomy alone can absorb a team for months), so I'd apply the principles, not install the framework:

- **Agree a simple taxonomy:** a handful of categories everyone accepts, not a six-month modelling exercise.
- **Use the data you already have:** the general ledger, the cloud bills, the data lake. Let AI do the heavy lifting of categorising it, rather than paying people to map it by hand.
- **Get to a view finance and technology both trust:** directionally right, not perfect.
- **Then run the real conversation:** where are we optimised, where aren't we, and what do we do about it.

The savings this surfaces are in run-rate spend you are already paying, so payback is fast, and the metric is simple: the share of technology spend categorised and visible. Wherever the group is standardising how it reads cost and margin, this is the technology view that belongs underneath it.

3. Gen AI skills distribution inside existing workflows

From the outside I can see many pockets of AI initiative across TTI. What I can't tell is whether they add up to an enterprise-wide effort to integrate Gen AI skills into workflows, get them evaluated and endorsed through the approval and risk controls you already run, and distribute them so the whole organisation gains. That distribution is reuse, which is what enterprise architecture is for. It's also one of the faster paybacks here: the hyperscalers and technology firms are enabling their people at a pace far ahead of most enterprises, and this is how that gap closes.

The measure follows the same rule as everything else in this document: each skill that gets endorsed and published carries a before-and-after measured on its own, so the returns are counted as they land, not promised up front. This initiative is aligned with my extensive background in service design.

The number this turns on

I won't hand you a benefit figure I haven't measured. Here is the hurdle instead:

At US\$299,563, this breaks even on a recurring saving of US\$132,500 a year: three-year benefit horizon, 10% discount rate, no benefit assumed in year one while the structure is built.

Recurring annual saving	3-year NPV	Return on fee
US\$132,500 (<i>breakeven</i>)	US\$0	1.0x
US\$500,000	US\$831,000	3.8x
US\$1,000,000	US\$1,961,000	7.5x
US\$2,000,000	US\$4,222,000	15.1x

Put the hurdle against the spend it governs. I don't have sight of your group technology spend. Whatever it is, US\$132,500 is the fraction of it this has to improve to wash its face. If the estate is anything like the scale I'd expect for a US\$15.3bn business, that fraction is small. But the first question is whether the measured estate actually clears the hurdle, and month one starts putting your own numbers on it.

Month 1 starts delivery, not a study

I'm not asking you to fund a month of discovery and then decide. Thirty days, three things delivered:

- 1. An architecture capability baseline.** What exists today, written down: the standards in force, where design assurance happens and where programmes self-certify, which forums hold which decisions and how long cross-region calls currently take to clear, and where the gaps are. The months 2-12 work builds from it.
- 2. Cost transparency, started.** The simple taxonomy drafted, and the first slice of spend categorised from data you already have: the GL and the cloud bills, with Gen AI doing the heavy lifting. Alongside it, the measurement of coordination and duplication cost defined and underway: I define what counts, so a renamed "alignment sync" doesn't quietly leave the baseline; your finance ops team runs the pulls.
- 3. Gen AI skills distribution, started.** The framework drafted: how a Gen AI skill (a repeatable way of doing a task) gets evaluated, endorsed, and published where every team can find it. The first candidate skill selected together, with its before-and-after measure defined: the template every skill after it reuses.

Every one of these is mine to deliver: none of them waits on anyone else moving first. Together they start all three capabilities visibly inside the first thirty days, turn the hurdle above into your actual numbers as the baseline fills in, and give you something concrete to hold me to for the remaining eleven months.

Where I've done this before. Each part of this has a precedent. The decision layer: at BG Group I built and chaired a global networks design authority spanning three regions: the standards, the exception process, and a board of regional design-authority peers deciding against one agreed target. Twenty-plus design standards came out of it, and design cycle time fell by up to 30%, measured. The assurance gate: as Head of Enterprise Solution Architecture

I restructured solution assurance across US\$60M+ of multi-tower solutions and lifted right-first-time delivery by 20%. The cost view: complex solutions ship with complex cost models, and I have built them throughout my career. The workflow side: twelve years leading service-design teams on end-to-end operational workflows.

Months 2-12 develop the capability from what you already have. This is not greenfield. Working from the month-one baseline of your current architecture capability, we define with the right stakeholders what the target capability should look like, complete the gap analysis between the two, and turn it into a candidate roadmap, then deliver it through a standard delivery framework, so the change lands as governed delivery rather than a side project. The result is cross-region decisions made once, by a defined owner, against a standard, with the cost view underneath. Not tactical fixes: a structure that holds after I've gone.

Why 12 and not 6. Six months buys tactical fixes that revert the moment I leave. Month one is baseline either way, governance takes three to four months to stand up, and you need two full quarterly cycles running through it to know the structure holds rather than the person holding it up. Six months stops before that test. The cheaper-looking option leaves less behind.

Why the work needs expert judgement. Gen AI makes candidate options cheap now: draft standards, reference models, documentation. What it doesn't do is judge which option is actually sound, which risk is real and which is noise, or when an exception is genuinely warranted. That discrimination comes from years of watching decisions go right and wrong, and it doesn't ship with the tool.

Why the role needs peer-level authority. The decisions this engagement settles cross director lines, so whoever holds them has to sit at that level to make them stick: a structural constraint, not a personal preference. Regional directors keep their teams, budgets and line authority throughout; the mechanism only settles the cross-region calls that already leave any one region's control.

Commercial. US\$299,563 all-in for the twelve months, full-time with normal leave. This is based on my previous Head of Enterprise Solution Architecture base salary of HK\$155,773 a month, with 25% added as there are no benefits, bonus or employer costs on top.

What this is not. Not a reorganisation and not a new function. Regional technology directors keep their teams, budgets and line authority. I work inside the initiatives you already have. The only decisions that reach a group layer are the cross-region ones already going to committee today.