

Enterprise Architecture & Technology Leadership

12-month engagement

Prepared for Ty Staviski, Deputy CFO, Techtronic Industries

Julian Hart · August 2026 · Private & confidential

At a glance: 12 months, US\$299,563 all-in, full-time, five capabilities in one engagement.

I'd work on the initiatives you already have, in an enterprise-architecture context. Not a parallel programme, not a new function competing for sponsorship.

The North American ERP conversion is proof of what's possible: a hard project done right because one team rallied around it. The risk with a win like that in a federated business is the lessons stay with the people who lived it and never reach the next region that needs them. Standards rarely fail at approval. They fail in adoption, quietly, region by region, until the reconciliation comes back.

You named the levers yourself: leverage synergies, AI and your assets. We'd build the governed layer that puts those levers to work, against your own target of 10% EBIT. With SG&A rising 173 bps in H1, we need to identify additional levers that ensure gross margin gains reach EBIT.

The next point could well come from the structural side. Steve asked exactly that at the annual results briefing: *"How can we leverage our negotiating costs on IT? How can we do the things globally to be able to free up more cash?"* Cost management frameworks like FinOps and TBM enable informed trade-off conversations between IT and finance. Architecture frameworks like TOGAF ensure adoption of global standards. It also feeds the free cash flow target raised mid-year. A region that reuses a group standard reduces capital up front, and robust assurance reduces capital spend on rework.

Where the cost reduction is: six levers

Lever	Mechanism
Capability mapping	In any federated business without a capability map, investment decisions default to the projects with the strongest sponsors, because there is no shared view of which capability each project serves or how it advances group priority. The map provides that view: every active technology project tied to the capability it serves, with a fund, defer or kill recommendation logged for each. The metric is portfolio coverage: the share of active projects mapped and carrying a logged recommendation. The investment decision stays with you; what changes is that it's made against the map, not the sponsor.
Portfolio rationalisation	Every application mapped to the capability it serves, with its run cost against it. Where two regions run separate platforms for the same capability (finance platforms are the classic case), each duplicate carries its own licences, support

contracts and upgrade cycle. Rationalising at contract end takes out run cost that recurs every year: a saving on money you are already spending, not a projected benefit.

Decision rights

Cross-region decisions escalate to committee because no one holds the right to make them. The governed layer gives each class of decision a named owner, a standard to decide against, a ratifier and a clock. The metric is time-to-decision: month one baselines how long cross-region calls take to clear today, and that number, tracked monthly, is what you measure me against. Every week a decision stops waiting is a week its benefit stream stops being deferred.

Design assurance

Programmes that self-certify their own design are how overruns and late re-cuts happen: the people accountable for the date are also the people judging whether the design is sound. An independent assurance gate separates those two jobs, so problems surface while they're still cheap to fix. The metric is coverage: the share of major programmes clearing independent design review before their dates lock. Overruns are your number, not mine. But rework falls where assurance holds, and the NA ERP conversion shows what that discipline looks like when it's in the room.

Vendor standardisation

Steve's point on the results call: leveraging negotiating costs on IT to free up cash. A reference architecture defines the standard platform per capability, so the estate converges from many regional variants toward few group standards. The metric is variant count per capability: baselined, with reduction candidates logged against contract end dates. Fewer variants means fewer, larger negotiations: the leverage Steve asked for. The negotiated saving is procurement's number; the variant count that makes it possible is mine.

Risk exposure reduction

A gap in the control estate is a cost waiting to land. Each additional cybersecurity control put in place, and the governance, risk and compliance processes around them, reduces the financial risk exposure the business carries. The metric is the controls and processes delivered, with the exposure reduction from each counted as it lands: a saving on losses you would otherwise be carrying as risk.

Scope: five capabilities, one engagement

All of it inside the enterprise-architecture context, working on the initiatives you already have. These are not five programmes. The first three are applications of the same architecture layer: standards, decision rights, cost visibility and reusable patterns. The last two are where else you can point me: deep security specialism, and senior delivery leadership on your priorities. And the layer works as a bridge between your US and APAC technology organisations: group objectives translated into executable plans in the regions, and regional realities surfaced back early enough to inform strategy rather than arriving as escalations.

1. The architecture capability

The design has to be group-level and the effort concentrates wherever the alignment gap actually is. Standing it up follows TOGAF's Preliminary Phase, tailored to TTI rather than adopted wholesale:

- **The capability itself:** which parts of the enterprise it covers, at what depth, and the target maturity: what you need now, evolving from there, not a full-dress framework on day one.
- **The operating model:** where the architecture function sits, the roles around it, and who is accountable.
- **Governance:** decision rights, approval and escalation paths, and how exceptions are handled.
- **Principles:** the enterprise-wide guardrails future architecture decisions get made against.
- **Fit with what you already run:** the frameworks, methods and delivery practices in place today; the capability is designed around them, not dropped on top of them.
- **Tools and skills:** the architecture repository, and the competency gaps that need closing for TTI's own teams to operate the capability.

The measure is the lever table above, made inspectable: portfolio coverage (the share of active projects mapped, each with a recommendation logged), time-to-decision on cross-region calls, assurance coverage on major programmes, and variant count per capability, all baselined in month one and tracked from there. The savings those numbers enable show up in your accounts; the numbers themselves are mine to move.

2. Cost transparency across the technology estate

You can't reduce a cost estate you can't read, and in most federated groups technology sits as one lump in the P&L. Two disciplines exist for this: TBM is the strategic, top-down version; FinOps is the real-time, operational one. Same principle: see cost clearly, tie it to value, decide together. A full TBM implementation is heavy machinery (agreeing the taxonomy alone can absorb a team for months), so I'd apply the principles, not install the framework:

- **Agree a simple taxonomy:** a handful of categories everyone accepts, not a six-month modelling exercise.
- **Use the data you already have:** the general ledger, the cloud bills, the data lake. Let AI do the heavy lifting of categorising it, rather than paying people to map it by hand.
- **Get to a view finance and technology both trust:** directionally right, not perfect.
- **Then run the real conversation:** where are we optimised, where aren't we, and what do we do about it.

The savings this surfaces are in run-rate spend you are already paying, so payback is fast, and the metric is simple: the share of technology spend categorised and visible. Wherever the group is standardising how it reads cost and margin, this is the technology view that belongs underneath it.

3. Gen AI skills distribution inside existing workflows

From the outside I can see many pockets of AI initiative across TTI. What I can't tell is whether they add up to an enterprise-wide effort to integrate Gen AI skills into workflows, get them evaluated and endorsed

through the approval and risk controls you already run, and distribute them so the whole organisation gains. That distribution is reuse, which is what enterprise architecture is for. It's also one of the faster paybacks here: the hyperscalers and technology firms are enabling their people at a pace far ahead of most enterprises, and this is how that gap closes.

The measure follows the same rule as everything else in this document: each skill that gets endorsed and published carries a before-and-after measured on its own, so the returns are counted as they land, not promised up front. This initiative is aligned with my extensive background in service design.

4. Cybersecurity

Deep background for me, not an adjacency, and it sits naturally inside the same governed layer. If security is where the pressure is, point me there:

- **Security architecture** across the key domains: network, cloud, identity, data and applications.
- **NIST CSF** applied across the enterprise security posture, from strategic architecture down to operational guidance. Much of the framework is process rather than technical controls, and that is exactly where internal teams run out of capacity.
- **GRC**: contributing to the formulation and execution of governance, risk and compliance frameworks, policies and practices across the technology environment.
- **A globally distributed manufacturing and supply-chain estate protected as one**, spanning Asia, the Americas and Europe, with particular focus on OT/IT convergence risk.
- **AI in security**, evaluated both ways: as a defensive capability, and as an evolving threat vector that needs structured governance.

The measure follows the same rule as the first three capabilities. Two numbers: work delivered as part of the NIST CSF programme, controls implemented, designs shipped and assured, and the financial risk exposure reduction from each additional control delivered.

5. Ad hoc projects

Where you or the team have a project that needs senior technical leadership, I'll take it on and land it, with a named outcome agreed before I start. Leading cross-functional delivery without a dedicated reporting structure is how I have always worked: rallying stakeholders, vendors and functional teams that don't report to me. For twelve months you have a full-time senior technology operator, and the priority list is yours.

The measure follows avoided-cost logic: days of senior technical leadership delivered, valued at the market cost of sourcing that leadership separately. Each project led is a hire you didn't make, and the capacity is already inside the fee.

Where the fee breaks even

I won't hand you a benefit figure I haven't measured. But the return has three sources, each measured in this document:

1. **Recurring savings from the architecture initiatives:** the levers above at work, duplicate platforms rationalised, vendors standardised, cost transparency acted on. Run-rate cost taken out of the estate and counted as it lands.
2. **Risk exposure reduction:** additional security controls delivered, each retiring measured financial exposure.
3. **Avoided leadership cost:** security architecture and senior technical leadership delivered from inside the fee instead of sourced separately.

An example for the first source, recurring savings: if the fee had to wash its face on those alone, it breaks even at US\$132,500 a year: three-year view, 10% discount rate, no benefit counted in year one while the structure is built. I don't have sight of your group technology spend, but if the estate is anything like the scale I'd expect for a US\$15.3bn business, that is a small fraction of it. In practice the return is more likely a blend of all three sources, in whatever mix reflects how you decide to use me.

Month 1 starts delivery, not a study

I'm not asking you to fund a month of discovery and then decide. But I'm also not going to promise all five capabilities started inside thirty days, because that is how nothing gets finished well. Here are the five candidate first deliveries, one per capability. In the first week we agree the priorities together; I deliver two or three of them inside the month:

1. **An architecture capability baseline.** What exists today, written down: the standards in force, where design assurance happens and where programmes self-certify, which forums hold which decisions and how long cross-region calls currently take to clear, and where the gaps are. The months 2-12 work builds from it.
2. **Cost transparency, started.** The simple taxonomy drafted, and the first slice of spend categorised from data you already have: the GL and the cloud bills, with Gen AI doing the heavy lifting. Alongside it, the measurement of coordination and duplication cost defined and underway: I define what counts, so a renamed "alignment sync" doesn't quietly leave the baseline; your finance ops team runs the pulls.
3. **Gen AI skills distribution, started.** The framework drafted: how a Gen AI skill (a repeatable way of doing a task) gets evaluated, endorsed, and published where every team can find it. The first candidate skill selected together, with its before-and-after measure defined: the template every skill after it reuses.
4. **The first cybersecurity delivery agreed.** I sit down with the people running the NIST CSF programme, understand where it is stretched, and agree the first piece of work I take on and deliver.
5. **An ad hoc project taken on.** Picked together in week one: something at the top of your agenda that is stubborn and needs an owner. I take it on and get it moving inside the month, so you see how I lead delivery before the structure work has even bedded in.

Of the five, the first three are mine to deliver alone; the last two take a steer from you and time with your teams, worth knowing when we pick. Whatever we agree is delivering by day thirty, giving you something concrete to hold me to for the remaining eleven months. And as the baselines fill in, the breakeven example above starts turning into your actual numbers.

Where I've done this before. Each part of this has a precedent. The decision layer: at BG Group I built and chaired a global networks design authority spanning three regions: the standards, the exception process, and a board of regional design-authority peers deciding against one agreed target. Twenty-plus design standards came out of it, and design cycle time fell by up to 30%, measured. The assurance gate: as Head of Enterprise Solution Architecture I restructured solution assurance across US\$60M+ of multi-tower solutions and lifted right-first-time delivery by 20%. The cost view: complex solutions ship with complex cost models, and I have built them throughout my career. The workflow side: twelve years leading service-design teams on end-to-end operational workflows. The security side: CISSP-certified, with a US\$25M+ managed-security portfolio built across my career, including solution lead on Telstra's largest ever cybersecurity engagement.

Months 2-12 develop the capability from what you already have, and keep delivering alongside it. This is not greenfield. Working from the month-one baseline of your current architecture capability, we define with the right stakeholders what the target capability should look like, complete the gap analysis between the two, and turn it into a candidate roadmap, then deliver it through a standard delivery framework, so the change lands as governed delivery rather than a side project. The architecture build is the spine, not the whole workload: the same months carry the other towers as you prioritise them, cybersecurity deliveries agreed with the NIST CSF programme, and the priority projects that need senior technical leadership, each run to a named outcome. The result is cross-region decisions made once, by a defined owner, against a standard, with the cost view underneath, and a year of delivery landed along the way. Not tactical fixes: a structure that holds after I've gone.

Why 12 and not 6. Six months buys tactical fixes that revert the moment I leave. Month one is baseline either way, governance takes three to four months to stand up, and you need two full quarterly cycles running through it to know the structure holds rather than the person holding it up. Six months stops before that test. The same arithmetic applies to the delivery side: a meaningful cybersecurity delivery or a stubborn cross-functional project runs a quarter or more, so a six-month engagement lands one of them, perhaps two, before the clock runs out. Twelve months turns that into a pipeline: initiatives picked, delivered and measured through the year, with the architecture capability compounding underneath them. The cheaper-looking option leaves less behind.

Why the work needs expert judgement. Gen AI makes candidate options cheap now: draft standards, reference models, documentation. What it doesn't do is judge which option is actually sound, which risk is real and which is noise, or when an exception is genuinely warranted. That discrimination comes from years of watching decisions go right and wrong, and it doesn't ship with the tool.

Why the role needs peer-level authority. The decisions this engagement settles cross director lines, so whoever holds them has to sit at that level to make them stick: a structural constraint, not a personal preference. Regional directors keep their teams, budgets and line authority throughout; the mechanism only settles the cross-region calls that already leave any one region's control.

Commercial. US\$299,563 all-in for the twelve months, full-time with normal leave. This is based on my previous Head of Enterprise Solution Architecture base salary of HK\$155,773 a month, with 25% added as there are no benefits, bonus or employer costs on top.

What this is not. Not a reorganisation and not a new function. Regional technology directors keep their teams, budgets and line authority. I work inside the initiatives you already have. The only decisions that reach a group layer are the cross-region ones already going to committee today. And it is not a team build: I arrive and deliver personally (the blueprints, the standards, the designs) and manage the stakeholder work myself.