

TTI - Ty Call Script

What this is, and why it exists

What it is: what Julian says on the Ty call - the opening line, five talking points, and prepared holding lines for the objections most likely to come. The written proposal Ty *reads* is a separate document, tti-ty-engagement-proposal.

Why it was created (9 Aug 2026): it lived inside the proposal file until now, two-thirds of the way down, behind the document Ty reads. That meant scrolling past the wrong content to reach the right content under time pressure, on the morning of a call that decides the engagement. Split out so the pre-call surface is one clean open.

How it is used: open this and nothing else before the call. send: NEVER - several lines here would be actively damaging if they reached Ty, Stephan or Tony Chung.

⚠️ THREE FILES MUST STAY IN SYNC (as of 9 Aug 2026)

File	Role	Authority
tti-ty-engagement-proposal	What Ty reads	Content
This file	What Julian says	Delivery
tti-ty-proposal-sources	The evidence behind every claim	Verification

They drifted apart once already, on 8 Aug, during the round-2 review fixes: the proposal was corrected but this file still carried superseded guidance, including a figure deliberately removed for being an overreach.

The rule: any edit to one requires checking the other two in the same session. Most at risk are the numbers (fee, hurdle, H1 figures), the month-1 design, and anything tagged INTERNAL / not-quotable in tti-ea-requirements.

★ THE LINE - say this early, close to these words

*"I want to spend the first weeks understanding the landscape and talking to the people running each line of business - but not as a study. I want to be delivering something you care about from week one. **Give me the thing that's stuck, that everyone knows needs doing and no one can get moving, and let me land it.** That's how I learn your organisation, and it's how you find out whether this works."*

Why this line and not the others (Julian, 8 Aug): he works best with something real to deliver against within days; a month of analysis is where he loses direction. So execution-first is not a sales device here, it is how the engagement has to be built to succeed. It also resolved the 7 Aug Codex FATAL - Ty cannot defend a no-break 12-month commitment where the value case only appears after signature. Both Codex rounds are now closed; the 8 Aug review's remaining fatal was verified as a false positive (the procurement quote is public, from the results call, not the internal town hall).

The frame underneath it, in Ty's head (from the Codex review):

"I'm not buying a theory of enterprise architecture. I'm buying a mechanism that takes one stuck cross-region technology decision and gives it an owner, a standard, a ratifier and a date."

△ **NEVER say "an extra pair of hands", "help out", or "where people are too busy".** The instinct behind those words is right; the words are a contractor framing. Stephan repeats Julian's phrasing verbatim, and it reaches Tony Chung as "we've got someone helping with the backlog" - which kills the director-level argument the whole engagement rests on. Say **"land it", "get it moving", "make it resolve"**. The early delivery is the *proof of the mechanism*, never a substitute for it.

The six talking points

1. **"What I'd like to do is work on the initiatives you already have, but in an enterprise-architecture context.**

Take the North American ERP win. Hard project, done right, because one team rallied around it. The risk with a win like that in a federated business is the lessons stay with the people who lived it and never reach the next region that needs them. You said it yourself: leverage synergies, AI and your assets. We'd build the governed layer that puts those levers to work." △ **Do NOT name the global P&L platform or the IT cost-leverage results** - both are internal (employee Q&A) and naming them reveals you hold an internal transcript. If you want to open that door, ask instead: *"Is there work underway to standardise how the business reads its cost and margin?"* See tti-ea-requirements § Quotability. △ **Do not say "AI agenda aimed at G&A"** - that overclaims. Ty's quote lists G&A cost reduction and leveraging synergies/AI/assets as separate drivers, not AI pointed specifically at G&A.

2. **"The reason that layer pays is that cross-region decisions currently have no owner, so they go to a committee.** Every one that waits defers whatever was behind it. The scope is five capabilities in one engagement: the architecture capability itself, cost transparency across the estate - so you can actually read what technology costs - Gen AI skills distribution: skills evaluated, endorsed, and published where every team can find them, with TTI's own teams operating it from there - and then two more ways to use me: cybersecurity, which is deep background, and senior technical leadership on whatever projects need it." **If he probes the cost discipline (reframed 9 Aug - supersedes the old "TBM, not FinOps / FinOps is the cloud corner" line, do NOT say that):** *"TBM is the strategic, top-down version; FinOps is the real-time, operational one. Same principle - see cost clearly, tie it to value, decide together."* Then the practical posture: principles, not a framework installation - a simple taxonomy everyone accepts, the data TTI already has (GL, cloud bills, data lake), AI doing the categorising, a view finance and technology both trust, directionally right rather than perfect. Fluency in how the two disciplines relate lands as senior; promising a full TBM programme is a rod for your own back. **On Gen AI, stay bounded** - the first skill through the pipeline, then their teams run it. Never sound like an AI transformation programme.

3. **"I'm not going to quote you a benefit number I haven't measured. But the return comes from three places.**

Recurring savings from the architecture work - duplicates rationalised, vendors standardised, run-rate cost out of the estate. Financial risk exposure retired by each security control delivered. And the leadership cost you avoid when I take on the security architecture or the technical lead myself. As an example, if it had to wash its face on recurring savings alone, it breaks even at about US\$132,500 a year on a three-year view at ten percent - but in practice it's more likely a blend of the three, in whatever mix reflects how you use me." **Don't add "it's only a few basis points"** - you don't know his denominator, and asserting it undercuts the measured-discipline posture that makes the rest credible. Let him do that arithmetic himself. △ **Reframed twice on 10 Aug (Julian's calls):** first the breakeven came off the proposal's at-a-glance line; then the section was rebuilt as the three-source return above, with the **NPV table removed** and US\$132,500 kept only as a savings-alone example, not a hurdle to be held to (the rod-for-own-back

fix: if big savings initiatives prove hard to find, the engagement earns its keep through exposure retired and leadership delivered). If Ty asks for return multiples, the old table re-runs in a minute from the arithmetic note in "Notes for Julian" - it is out of the written document by design, not by omission.

4. **"Month one starts delivery, not a study.** There are five candidate first deliveries, one per capability - and I won't pretend I can start all five in thirty days; that's how nothing gets finished well. In week one we agree the priorities together, and I deliver two or three of them inside the month. The candidates: a baseline of the architecture capability you have today: the standards in force, where assurance happens and where programmes self-certify, which forums hold which decisions and how long cross-region calls take to clear. The first cut of the cost view: the taxonomy drafted and a first slice of spend categorised from your own GL and cloud bills. The first cut of the Gen AI skills-distribution framework - how a skill gets evaluated, endorsed and published - with the first candidate skill chosen together. The first cybersecurity delivery agreed: I sit with the people running the NIST CSF programme, find where they're stretched, and agree what I take on first. And an ad hoc project, picked together in week one - the stubborn one at the top of your agenda that needs an owner. I take it on and get it moving." **△ Never say "quick win"** - Kari sold the 12 months on strategic gains, not quick wins. **△ SUPERSEDED 10 Aug: the stuck-project play is WRITTEN again** (reverses the 9 Aug "unwritten overdelivery" ruling, Julian's call: month-1 item 5 now commits to taking on a top-of-agenda stubborn project, and item 4 commits to agreeing the first NIST CSF delivery - **THE LINE** and the written list now match). The guardrails that survive unchanged: **never promise anything RESOLVES inside 30 days** - the written commitment is take-on-and-get-moving; and on a stuck cross-region *decision*, the decision is still made by the owner Ty names. Speakable: *"I don't make the decision - your named owner and ratifier do. I build the machinery that gets it made."* **Items 1-3 are artefacts Julian ships alone; items 4-5 take Ty's steer and team time** - that split is stated honestly in the document itself, so don't oversell sole control. **And never promise all five in month one (Julian, 10 Aug: "try and do all five, and I'm going to fail")** - the written commitment is two or three, priorities agreed in week one. If Ty pushes for all five, the answer is prioritisation, not capitulation: *"we can start them all, but two or three delivered beats five half-done - which matter most to you?"*

5. **"It's twelve months, not six. Six buys tactical fixes that revert when I leave, it lands one initiative, maybe two, before the clock runs out, and it's actually dearer per month."** (delivery-pipeline clause added 10 Aug to match the widened "Why 12 and not 6" in the proposal: a real cyber delivery or a stubborn project runs a quarter or more, so twelve months buys a pipeline of them, not one.) Then the number, stated plainly:

*"It's **US\$299,563 all-in for the twelve months**, full-time with normal leave. That's based on my previous **Head of Enterprise Solution Architecture** base salary of HK\$155,773 a month, with 25% added as there are no benefits, bonus or employer costs on top."*

△ Say "Head", not "Global Head" (Julian, 9 Aug): "Global" invites the counter that this role isn't global scope, so the comp basis gets argued down. The verified full title stays Global Head if accuracy is ever challenged - understating is safe.

△ Do not mention benchmarking. The goal is that he signs it off, not that he goes and compares. The base figure is there because it explains where the number comes from - and it is available to him if he *is* under pressure to benchmark - but framing it as a comparison aid invites the exercise Julian just lowered the rate to avoid.

Say the role, never "market rate". The role tells Ty what level the number bought and is an unarguable fact about Julian's history. **A market-rate claim is contestable and the evidence is genuinely split** - see the holding line below.

Do NOT say "including the pension contribution my employer made" - that was the June wording for the higher figure, and quoting it now would be inconsistent with the number being charged. **Do NOT open on HK\$194,716** - it

is the arithmetic result, not the framing.

6. **The versatility close** (added 10 Aug, from Stephan's 10 Aug steer: "tell him that you are happy to work with the team and him and any other ad hoc projects"): "And beyond that scope, use me where the need is. Cybersecurity is deep background for me - CISSP, security architecture, a US\$25M managed-security portfolio built - and if you or the team have projects that need senior technical leadership, point me at them. I'll land them." **Spoken, never say "extra pair of hands", "help out", or "ad hoc" aloud** - Stephan's phrase, but it is contractor framing and Stephan repeats Julian's words verbatim (the standing guardrail at the top of this file). The written proposal carries "Ad hoc projects" once, as the capability-5 heading, framed at senior altitude ("senior technical leadership... named outcome agreed before I start"); spoken, it is always "projects that need senior technical leadership". **Restructured later same morning at Julian's direction: cyber and ad hoc are now capabilities 4 and 5 of a five-capability scope**, so talking point 2 above says five, and this close is the elaboration of capabilities 4-5, not a separate section. **Capability 4 was then rebuilt as a work description** (five bullets from the April role profile: domain sweep, NIST CSF posture work, GRC, OT/IT manufacturing protection, AI-in-security) with the credentials (CISSP, US\$25M portfolio, Telstra) moved to "Where I've done this before" - so if Ty probes cyber, speak the WORK first and let the credentials answer "where have you done this"; don't lead with the CV. On AI-in-security, hold the boundary: it is one bullet inside capability 4, evaluated-and-advised - never merge it with capability 3's skills-distribution pipeline, and if someone says "we already have AI governance" the existing holding line applies. **Why this exists (internal)**: the ultimate goal is full-time employment, so versatility - going where the need is, including BAU, cyber, or technical-lead work - is the real sell, and it stops the whole case hanging on finding US\$132.5k/yr savings initiatives.

(Role-profile mapping pass, 10 Aug, later same morning - three more lines now in the written proposal, speakable if the moment fits:) the **US-APAC bridge** ("the layer works as a bridge between your US and APAC organisations - group objectives translated into executable plans in the regions, regional realities surfaced back early enough to inform strategy rather than arriving as escalations" - this is Ty's own 15 Jul framing handed back to him); the **practitioner line** in What-this-is-not ("not a team build - I arrive and deliver personally"); and **delivery without line authority** in the ad hoc bullet ("rallying stakeholders, vendors and functional teams that don't report to me"). The cyber bullet now sweeps five named domains: network, cloud, identity, data, applications.

Holding lines

If he pushes on the number: "This is what I cost my last employer, and I now carry my own severance, notice and cover." Do not go up. Do not offer a range. **And do not go down** - the rate was already lowered on 9 Aug from US\$320,644 specifically to stay close to Tony's band, and at US\$299,563 there is HK\$4,421 of headroom, effectively nothing. Any further concession is below cost recovery.

If he asks for six months: "I'd rather not - you'd be paying more per month for work that reverts." (Reason if probed, spoken-only - removed from the written proposal 9 Aug to keep Ty's attention off Julian's cost structure: the annual fixed costs Julian now carries himself don't halve over a shorter term, so a 6-month rate is higher per month - HK\$218,252 vs ~HK\$194,716. Full basis: independent-consulting-pricing.)

If he asks "and what if month one shows the estate doesn't clear the hurdle?" (the no-off-ramp question - the sharpest unanswered objection from three review rounds; line adopted 9 Aug, do not improvise past it):

"Then you'll see it in writing in month one, with the numbers - not discovered in month eleven. Month one isn't a go/no-go gate; there's no break clause and I won't pretend it works as one. What it does is size and aim the months 2-12 plan. I've priced this on my judgement that for a business your size the addressable pool is a multiple of US\$132,500 - and month one replaces my judgement with your numbers."

Why this framing: it concedes the structure honestly (no gate, no break clause) instead of letting Ty discover the inconsistency himself; it converts the fear into a transparency promise with a date on it; and it restates the confidence claim as one Julian is accountable for in writing at day 30. Do not offer a break clause to soften it - that decision is locked. (10 Aug note: this objection is now structurally weaker - the written proposal no longer stakes the fee on the savings hurdle alone but on a blended return across savings, exposure retired and leadership delivered, so a thin savings pool no longer means the engagement failed its own test. The holding line above still works as written; the blend is the additional answer if he pushes.)

If he challenges the level - "is that what someone at that level actually earns here?"

"Robert Half's Hong Kong guide puts enterprise architect at HK\$2.16M at the seventy-fifth percentile - and they define that percentile as extensive experience, advanced skills, specialised certifications. That's the profile: twenty-five years, TOGAF, CISSP, CCIE. My base was HK\$1.87M. So I was below the level their own guide says that profile commands."

Verified live 9 Aug - figures and definitions confirmed on the Robert Half site (hk-architecture-salary-benchmarks §2).

⚠ **Use only if directly challenged, and only this source.** Randstad's guide caps enterprise architect at HK\$120,000/month, which would put Julian 30% above their ceiling. The two verified sources disagree by 50% and the gap is not resolvable from published data. **Volunteering a comparison hands whoever is checking the choice of which guide to use** - and Randstad is the one that hurts.

Never claim "market rate." The one band that would cleanly support it (Head of EA at HK\$1.8-2.2M) has no primary source and is quarantined in Julian's own research as an AI inference, not market data.

If he raises overruns, or you sense the history is in the room: *"The reason those programmes overshoot is usually that the people accountable for the date are also the people judging whether the design is sound. Assurance separates those two jobs - nobody marks their own homework."* **Never name a programme, a number or a business unit.** The CEO's public position is that the North American ERP executed flawlessly, and it is Shane Mall's business. Let Ty supply his own example. Background: tti-ea-requirements § The overrun question.

If he asks which decision you'd take on: you don't have one, and that is now the stated position in the proposal rather than a fallback - you pick it together in week one against the three criteria (crosses regions, senior sponsor waiting, nameable cost of delay). Say it as design, not as a gap. **Do not invent one.**

If he asks "where have you done this before?" (the written proposal now maps a precedent per capability; spoken, still lead with the one clean BG Group story): *"The closest precedent is BG Group. I built and chaired a global networks design authority across three regions - the standards, the exception process, a board of regional design-authority peers all deciding against one target. Twenty-odd design standards, and design cycle time down up to thirty percent, measured. What I haven't held is a charter to retire business steering committees at group level, which is exactly why month one puts inspectable artefacts on your desk rather than asking you to trust a track record."* **Say "measured, not estimated" if he probes the 30%.** The honest-gap sentence is spoken-only - removed from the written proposal 9 Aug. Full story: tti-story-bank §1a.

If he goes deeper on Gen AI - "what would you actually do?" (added 9 Aug; reframed same day from "governance" to skills distribution): *"The contribution is bounded. I'd set up the way a Gen AI skill - a repeatable way of doing a task - gets evaluated, endorsed through the approval and risk controls you already run, and published where every team can find it. Repeatable, checked, endorsed, with defined levels of human judgment - and TTI's teams operate it from there. The point is the first savings stand up to audit, so the ones after them scale. Not tools*

built in a corner, and not an AI transformation programme." This is the only surface carrying the "not a transformation programme" guardrail - the proposal no longer says it, so it has to be said here if the AI conversation opens up. △ **Deliberately NOT called "governance" (Julian, 9 Aug):** AI governance is the security team's turf and claiming it invites a turf fight; this is the distribution and enablement layer. If someone says "we already have AI governance," the answer is: *"Good - that stays exactly where it is. This is the layer that gets endorsed skills into people's hands."*

If he asks where you've done the cost-transparency work before (added 9 Aug): *"The cost-view build itself I've done over and over - complex solutions always come with complex cost calculators, so pulling cost data into a structured view finance can interrogate is bread and butter for me. What I haven't done is get a cost taxonomy agreed across a whole enterprise - and that's exactly why the design is a handful of categories, your existing data and a directionally-right view, not a six-month modelling exercise. I'm applying the principles that make these programmes work and deliberately avoiding the machinery that makes them stall."* Same honesty device as the BG Group line: name the gap, then show the design already answers it. **Do not overclaim TBM programme experience** - Ty may know TBM well, and the credibility spine of the whole proposal is not claiming what hasn't been done.

If he presses on whether one example is enough (updated 9 Aug - *the written proposal now maps a precedent per capability, so the breadth case is on the page*): the second example - Global Head of Enterprise Solutions Architecture, US\$60M+ of multi-tower solutions assured, **+20% right-first-time** from restructuring solution assurance - is now IN the proposal, not in reserve. Spoken, still lead with BG Group alone (one clean story beats two competing); if he presses, walk the proposal's mapping: decision layer BG Group, assurance gate the +20%, cost view the cost models, workflow side the twelve years. Different engagements, same mechanisms. **The honest-gap line is now spoken-only** (removed from the written proposal 9 Aug, Julian's call): on the page it planted doubt; said aloud under questioning it builds trust. It lives in the "where have you done this before" answer above - use it there, don't reintroduce it in writing.

If he asks how he'd measure me - "what do I hold you to?" (added 9 Aug from *perplexity-defensible-ea-benefits; the leading/lagging split*)

"You hold me to numbers I control, and you can inspect them at quarter-end: how much of the project portfolio is mapped with a recommendation logged, time-to-decision on cross-region calls, assurance coverage on major programmes, variant count per capability. The savings those enable show up in your numbers, not mine - I won't claim them, because the moment an architect starts claiming other people's savings you should stop believing him. What I will show you, every quarter, is the metric and the decision it fed."

The four metrics match the lever table in the proposal (Capability mapping, Decision rights, Design assurance, Vendor standardisation rows) - if either surface changes its metrics, change the other. **They are also now gathered into capability 1's closing measure paragraph** (added 10 Aug), so this holding line and that paragraph must stay word-consistent: same four metrics, same "savings show up in your accounts, the numbers are mine to move" split. **The lever table gained a sixth row on 10 Aug (Risk exposure reduction: cyber controls + GRC processes delivered, exposure reduction counted as it lands)** - that row's metric belongs to capability 4's measure, not capability 1's four, so don't fold it into this answer unless the cyber capability is live in the conversation. **If capability 4 (cyber) activates, its measure is in the proposal too** (added 10 Aug): work delivered as part of the NIST CSF programme (Julian's call, 10 Aug: the "or whichever programme you run" hedge removed - TTI's use of NIST CSF is publicly evidenced, see tti-cybersecurity, so the direct reference is safe and lands as knowing the environment), plus financial risk exposure reduction per control. **The mechanics are spoken-only (trimmed from the written version 10 Aug, Julian's call - the written measure stops at "each additional control delivered"):** if Ty probes how exposure reduction is counted, the answer is exposure before versus after (impact times probability), counted in a

model TTI's security leadership owns, not Julian's - same never-claim-a-number-from-a-model-he-doesn't-own discipline as everywhere else. **Capability 5 (ad hoc) also carries a measure (added 10 Aug):** avoided-cost logic - days of senior technical leadership delivered, valued at the market cost of sourcing it separately; each project led is a hire not made, and the capacity is already inside the fee. Speakable: *"every project I lead is senior leadership you didn't have to buy elsewhere - it's already in the fee."*

Why this framing: it is the same discipline as talking point 3 - refusing to claim what isn't measured or owned - applied to his accountability question instead of the benefit question. Do not soften it into shared credit; the strength is the refusal.

If he engages on the H1 numbers: with SG&A rising 173 bps in H1, you need additional levers to make sure gross margin gains reach EBIT. Gross margin expanded 258 bps; only 86 of it got there. The next margin point could easily come from the structural side. Steve asked exactly that at the annual results briefing earlier this year: *"How can we leverage our negotiating costs on IT? How can we do the things globally to be able to free up more cash?"* Cost management frameworks like FinOps and TBM enable informed trade-off conversations between IT and finance. Architecture frameworks like TOGAF ensure adoption of global standards. **△ Two things to hold separately:** attribute this to Steve, not Ty - it's a different quote from Ty's structural G&A line. **And it's from a different, earlier event than the H1 numbers** - the annual results call (FY2025) versus the H1 2026 interim announcement. Don't imply they were said together. The free cash flow target was also raised mid-year - he knows the numbers, no need to restate them. Avoided technology capex **supports** that too, two ways: a region that reuses a group standard reduces capital up front, and robust assurance reduces capital spend on rework. **△ Don't say "you raised the target"** - this is the Board's collective announcement, not something to attribute to Ty personally. He's not even a named director on this filing; Frank still is. **Say "supports", not "flows straight into"** - replacement opex and timing complicate it, and overstating a cash-flow mechanism to a CFO is the one place precision matters. Full workup: tti-h1-2026-results-talking-points.

Notes for Julian (do not send)

- **Calibration - what this engagement actually is (Julian, 9 Aug).** This is sponsored by Stephan and is, in reality, the route into the company: they know Julian, they know they're getting the equivalent of a full-time employee for 12 months, and nobody on their side is modelling a take-the-money-and-walk scenario. So: play the game properly - demonstrate the value, keep every commitment an artefact within Julian's control - but don't overharden the document against an adversarial-procurement reading that isn't the real audience. There is likely leniency around the no-break-clause in practice. This calibration is context for preparation only; it changes nothing about what is said on the call.
- **The framing risk is live.** Stephan repeats Julian's words verbatim to Ty. Anchor on *establishing architecture and governance across the group* (updated 9 Aug from "across Asia": the proposal is group-level scope, and an Asia-only framing relays a smaller engagement than the one priced); never "enterprise architect" as a title, never "IT support", never "consulting gig".
- **Keep the fee off anything that travels.** If Ty asks for a version to circulate to procurement, strip the Commercial section first. Rationale in independent-consulting-pricing.
- **Tony Chung sequencing.** He must hear about this before it reaches a sign-off desk, framed as pooling authority rather than confiscating it. The "What this is not" paragraph in the proposal exists for him, not for Ty.
- **The NPV arithmetic is checkable:** $\text{US\$}299,563 \div 2.2607 = \text{US\$}132,505$. The 2.2607 is the 3-year annuity factor at 10% with benefits falling in years 2, 3 and 4 ($0.8264 + 0.7513 + 0.6830$). If Ty supplies TTI's WACC the whole table re-runs. (10 Aug: the NPV table itself is out of the written proposal - the breakeven survives only as a

savings-alone example. This arithmetic stays here to defend the example and re-create the table on request: saving $\times$ 2.2607 – fee.)

- **The 10% discount rate, explained (added 9 Aug, for if Ty asks "why ten percent?").** A saving that arrives in two years is worth less than the same saving today, because the money tied up meanwhile has a cost - the business could be earning a return on it elsewhere. The discount rate is the annual rate used to shrink future savings back to today's value. At 10%: a saving landing in year 2 counts at 82.6 cents on the dollar ($1 \div 1.10^2$), year 3 at 75.1 ($1 \div 1.10^3$), year 4 at 68.3 ($1 \div 1.10^4$). Those three factors are what sum to the 2.2607 annuity factor in the arithmetic note above - benefits are assumed to start only in year 2, because year 1 builds the structure. **Why 10% specifically:** it is the conventional, conservative round-number stand-in for a company's cost of capital (WACC) when you don't know their actual figure. Direction to hold under questioning: a **higher** rate shrinks future savings more, so the breakeven saving **risks**; a lower rate does the opposite - so 10% being roughly right or slightly high is the conservative position. **The proposal deliberately makes no offer to re-run at TTI's WACC** (the "give me your WACC" line was removed 9 Aug - if Ty wants it at his rate he can ask, or run it himself; it's his prerogative, not a service to volunteer). If he does ask, the table re-runs in a minute from the arithmetic note above. **Speakable version:** "Ten percent is a conservative stand-in for your cost of capital - money later is worth less than money now, so I discount the future savings back to today's value before comparing them to the fee."
- **TTI IT spend is deliberately not asserted.** It is not in the vault and guessing it in front of a Deputy CFO loses the room. The proposal hands him the hurdle and asks him to supply his own denominator - the same device as letting him set the multiple (tti-ty-call-simulation-2026-07-19).
- **The procurement line IS safe to say.** Steve's "how can we leverage our negotiating costs on IT" is from the **public** annual results call (tti-2025-annual-results-script line 435), logged as P3 in tti-ea-requirements. **Do not confuse it with T12** - a similar-sounding IT-cost-leverage line about Ty's organisation that came from the employee Q&A and is internal/not quotable. A Codex review conflated the two on 8 Aug; the distinction is real and worth holding.
- **How month 1 "settles the hurdle question", mechanically (added 9 Aug - so the claim can be defended step by step).** Month 1 measures the **addressable pool**, not savings. (1) Julian defines what counts: the coordination tax (committee and escalation time on cross-region decisions, from calendar data, hours $\times$ loaded rates - the definition rule is why a renamed "alignment sync" can't leave the baseline) and the duplication cost (run cost of overlapping applications, from portfolio and contract data). (2) TTI's own finance ops pulls the data: the baseline comes from Ty's numbers, counted by his people, against rules Julian sets - neither side marks its own homework. (3) The output is an annual recurring figure, the pool, and the hurdle question becomes "is capturing this fraction of the measured pool credible?" - a division Ty does himself, the same device as fee-versus-spend. (4) The months 2-12 planning (target state, gap analysis, roadmap - the ADM cycle in the proposal's "Months 2-12" section) is then priced against that baseline as it lands: pool converted to capture, lever by lever. **The plan is no longer a month-1 deliverable** (cut 9 Aug - no plan before target state and gap analysis exist). (5) If the pool measures small, Julian says so in writing - the no-unmeasured-numbers spine applied to his own case. "Settled" means the benefit case has been converted from Julian's assumption into TTI's own measured numbers, either way. **Timing honesty (revised 9 Aug): month 1 commits to the measurement being defined and underway, not complete** - calendar and portfolio pulls in a big org can run past day 30 through nobody's fault, so the pool number lands "as the first numbers land," possibly early month 2. The written promise is the floor; finishing inside 30 days is the overdelivery case.
- **The cost-transparency closing passage, decoded (added 9 Aug - what each sentence is doing, so it can be defended fluently).** (1) "Run-rate spend you are already paying, so payback is fast" - run-rate is the recurring money that leaves automatically every year (licences, support, cloud, renewals). Savings found there mean stopping a payment already flowing, so they start immediately and repeat - no benefit lag, no dependency on

behaviour change, unlike a transformation benefit. (2) *"The share of technology spend categorised and visible"* - the metric Ty holds Julian to: what percentage of TTI's technology spend is mapped into the taxonomy. Near zero today (one lump in the P&L), rises as mapping progresses, and entirely architect-controlled - it measures whether the view got built, not whether anyone acted on it. (3) *"Wherever the group is standardising how it reads cost and margin..."* - the deliberate conditional. Julian knows internally (employee Q&A, NOT quotable) that a group cost/margin standardisation initiative exists; the "wherever" phrasing asserts no knowledge, and invites Ty to make the connection himself - repositioning the capability as the missing technology layer of something already underway. Spoken version of the same device: *"Is there work underway to standardise how the business reads its cost and margin?"*

- **Taxonomy-risk guardrails (9 Aug, from the TBM evidence review).** Julian's two stated fears - a taxonomy that never gets agreed, and a taxonomy nobody uses - are the two failure modes the published evidence documents. The success pattern is to attach the taxonomy to an allocation model, a named decision-owner, a savings target and a review cadence; without those it becomes "an artifact of governance aspiration" (the GAO's federal TBM audit found transparency but no identified savings for exactly this reason). The four-step design in the proposal already encodes the mitigations, and month 1's measured baseline is the decision-attachment. **Never quote the case-study numbers to Ty** (Bank of Ireland, MassMutual, National Grid, SBA, Washington State) - vendor-published, non-TTI analogues, same ruling as the ISG and Thai Union figures. The "15-30% of governed budget" governance-cost claim is unverified - do not repeat it anywhere.
- **FinOps payback speed is stated qualitatively, not with a benchmark percentage.** Deliberate: the proposal rests on "I won't quote what I haven't measured", and a vendor statistic would break that discipline at the one point a CFO would check it.
- **Pick two concrete examples before the call (round-3 S7, confirmed 9 Aug):** one real cost-model example (a specific complex solution whose cost calculator Julian built - name the deal shape, not the client if sensitive) and one service-design workflow example (a specific end-to-end workflow his team designed). The written claims are deliberately broad; these are the show-me answers behind them. Candidates in tti-story-bank.
- **Ty reads the document before the call (decided 9-10 Aug, per Stephan's steer: send the PDF beforehand).** So the call may open with Ty's reaction or question rather than Julian leading with THE LINE. Every talking point above must work as an answer, not only as an opener - THE LINE still gets said, but it may land second, not first.
- **Still open before the call: the right-to-work question - now unresolvable on this timeline: Kari must NOT be contacted (his mother died, per Stephan 10 Aug).** The decision-speed question is ANSWERED (Stephan, 10 Aug: "Ty should be able to make a pretty quick decision"; he has already reminded Ty about the call).

Related

- tti-ty-engagement-proposal - **the paired document.** What Ty reads; authoritative on content. Keep in sync.
- tti-story-bank - §1a the BG Group prior-proof story, §3 predictable role questions
- tti-ea-requirements - the Quotability table: what can and cannot be said out loud
- tti-h1-2026-results-talking-points - the H1 2026 financial detail behind the numbers above
- tti-ea-governance-value - the full argument bank, if a line needs its underlying reasoning